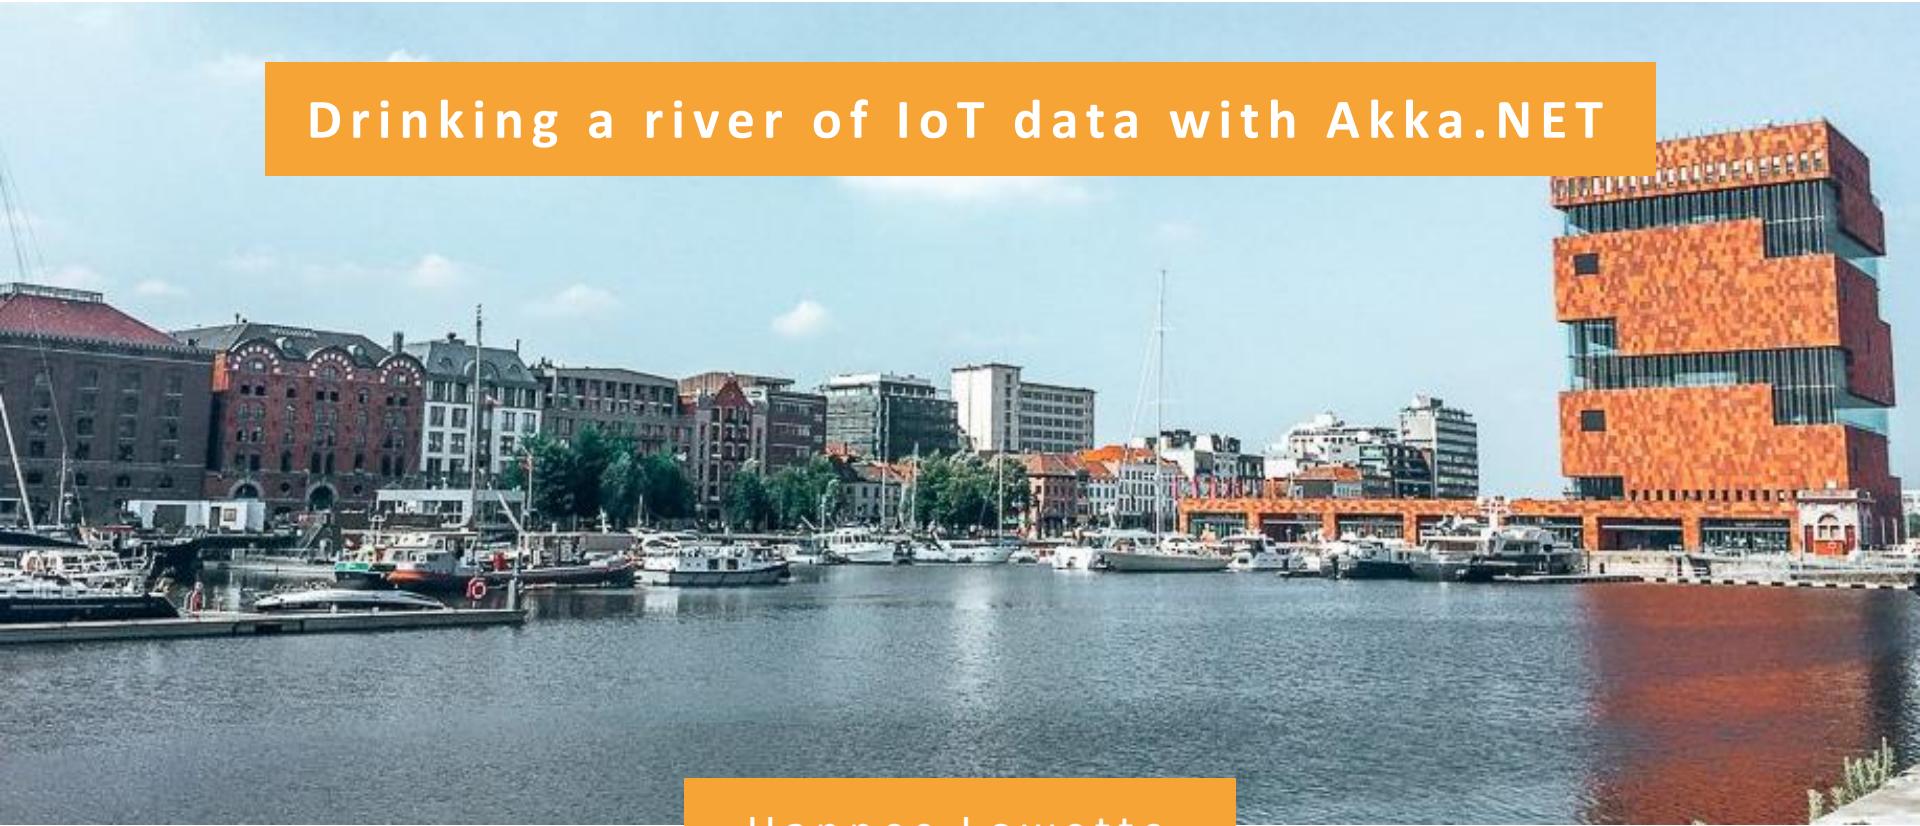


Drinking a river of IoT data with Akka.NET



Hannes Lowette



- 1 Introduction to Akka.NET
- 2 The problem domain
- 3 How does Akka.NET fit?
- 4 Implementation details
- 5 Beyond this talk ...

Introduction to Akka.NET

History & Principles

IT IS ABOUT PEOPLE



Actor model: the origin

Research paper in 1973:

- Carl Hewitt, Peter Bishop and Richard Steiger
- Designing software inspired by physics
- Aimed at many independent microprocessors

Further refinement:

- Irene Greif → operational semantics for the Actor model
- Henry Baker & Carl Hewitt → axiomatic laws for Actor systems
- William Clinger → denotational semantics based on power domains (1981)
- Gul Agha → transition-based semantic model (1985)



1980's: Major adoption

Ericsson AXD 301 Telco System:

- Invention of Erlang
- Fault-Tolerant
- Distributed
- Concurrent
- 2 million lines of code
- 99,999999% uptime (9 nines)

(31ms downtime per year)



ERICSSON

Evolution of Akka.NET

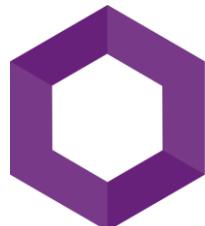
- **2010:** Akka on JVM (Scala)
- **2013:**
 - Aaron Stannard: started porting Akka to .NET
 - Roger Johansson: started porting Akka to .NET (Pigeon)
- **2014:**
 - Combined codebases & effort
 - Permission from Lightbend to call it Akka.NET
- **2015:** First version of Akka.NET released
- **Today:** team expanded, more stable & complete, ...



What happened in 2015?

3 Actor frameworks for .NET get released:

- **Feb 2015:** Project Orleans v 1.0.0
- **April 2015:** Akka.NET v 1.0.0
- **April 2015:** Service Fabric Reliable Actors v1.0.x



Orleans



Microsoft Azure
Service Fabric



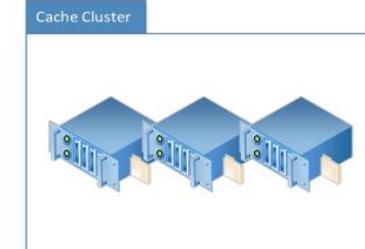
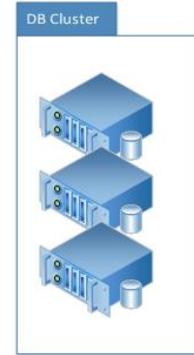
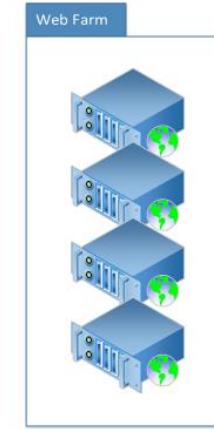
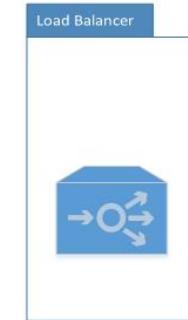
why 2015?

If these concepts exist since the 70's...

Classic scaling can't keep up

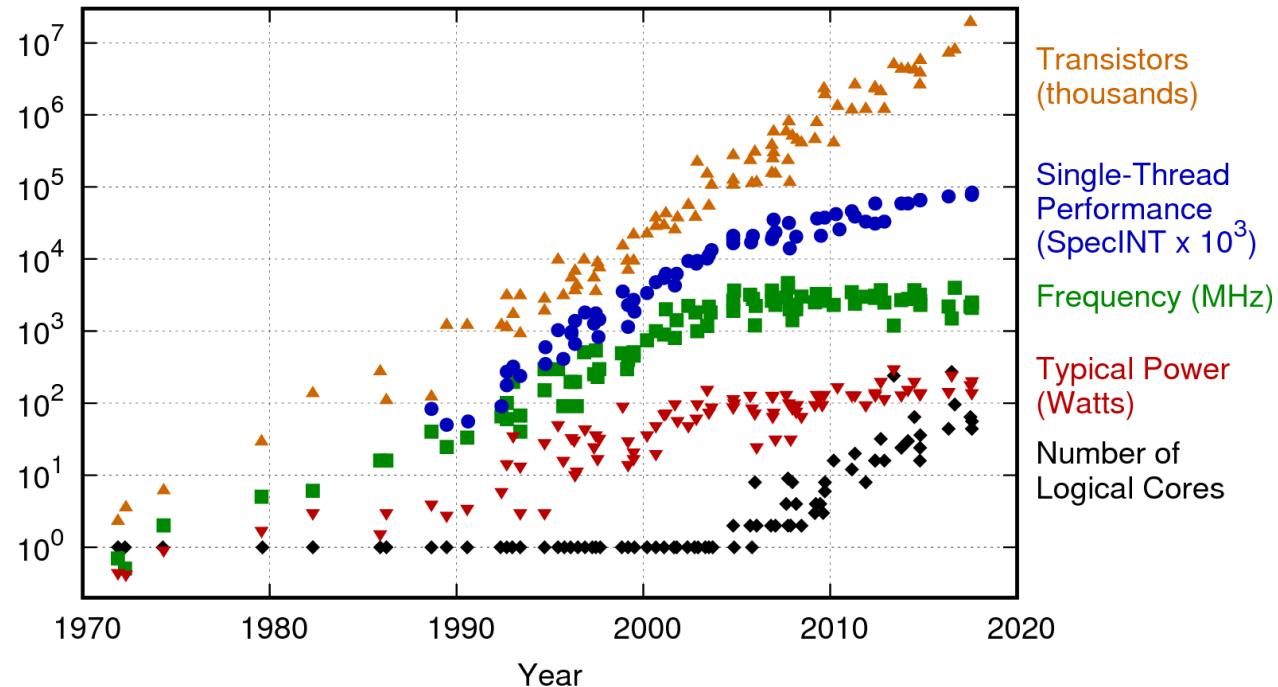
Explosion of # consumers:

- Smartphone apps
- Internet of things
- Regular website traffic



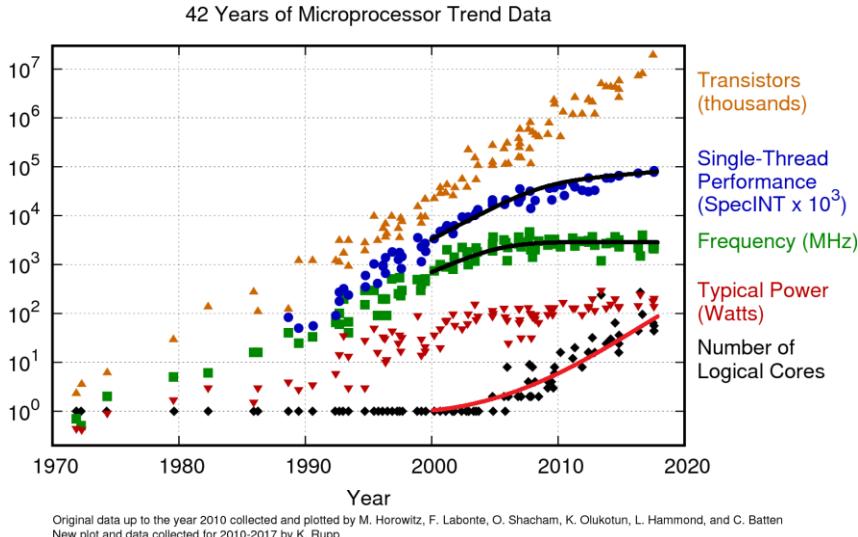
Free lunch was over

42 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

Free lunch was over



- Clock frequency has stabilized for over a decade
- Single thread performance growth is slowing down

But ...

The number of cores is increasing!

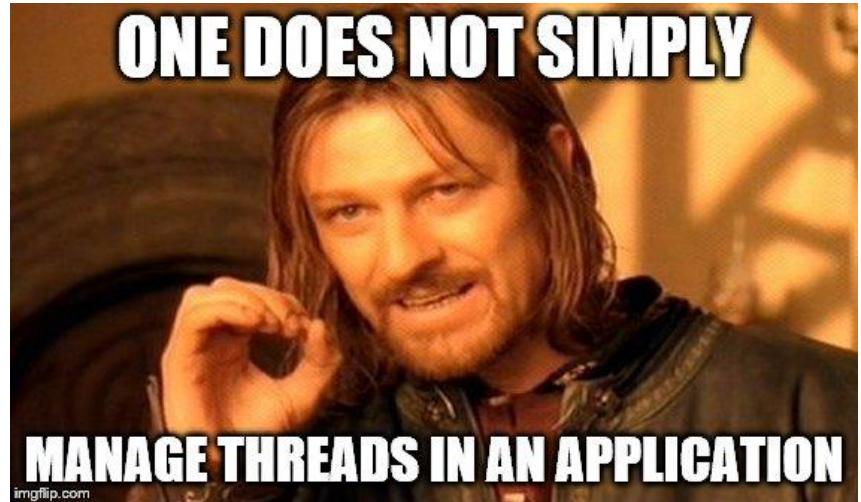
Parallelism is the salvation

Problems of parallel execution:

① Shared state:

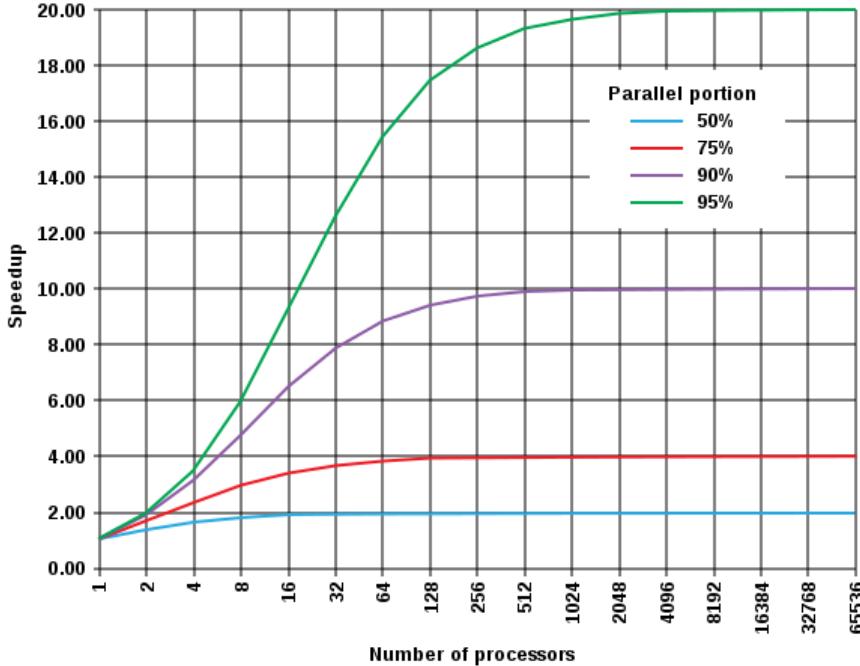
- Race conditions
- Blocking calls
- Deadlocks

① Some code can't be parallelized



Serialization hurts scalability

Amdahl's Law



- The theoretical speedup
- For a fixed workload
- When increasing resources
- Depending on how much can be parallelized

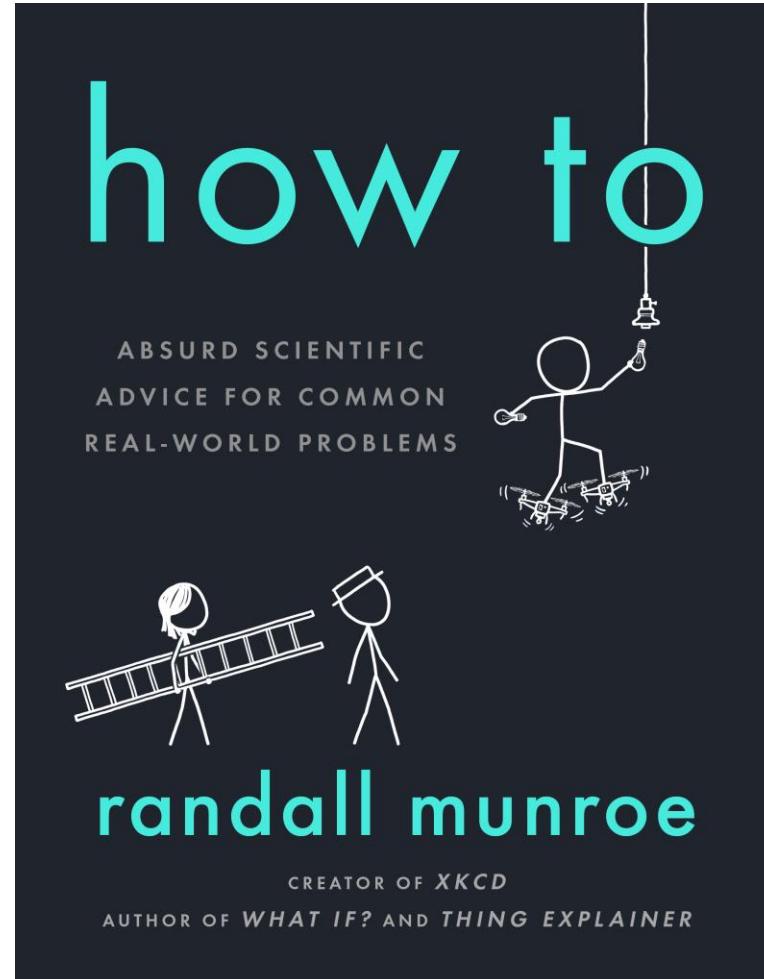
90% isn't going to cut it.

Can actor models help?

Actor models bring:

- High degree of parallelization
for stateful systems
- Reactive patterns
- Fault tolerance (self healing)

But, HOW?

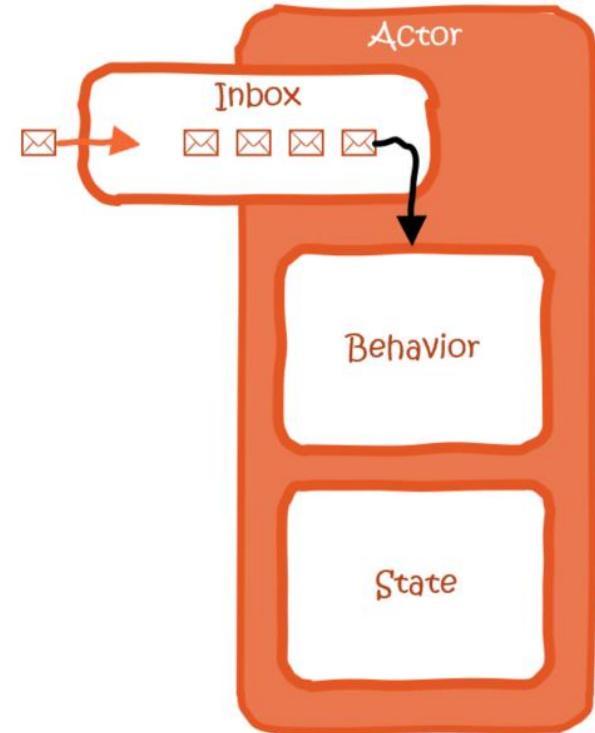


The Actor

Simple object

- Holds its own state (no shared state)
- Only input: messages
- Inbox:
 - Processed in order
 - 1 message at a time

→ An Actor runs single threaded
→ no concurrency problems



The Actor

```
public class MyActor : UntypedActor
{
    protected override void OnReceive(object message)
    {
        if (message is MyMessage myMessage)
            DoSomething(myMessage);
    }

    private void DoSomething(MyMessage myMessage)
    {
        // TODO: handle the message here
    }
}
```

Messages

- Simple objects
- Immutable!
 - Akka.NET doesn't enforce this
 - You **SHOULD NOT** try to exploit this
- Might cross machine boundaries
- Throughput:
 - Claimed: 50 M/s on a single machine
 - Well over 1 M/s on my laptop



Messages

```
public class MyMessage
{
    public int IntProperty { get; }
    public string StringProperty { get; }
    public ImmutableArray<decimal> Values { get; }

    public MyMessage(int intProperty, string stringProperty, ImmutableArray<decimal> values)
    {
        IntProperty = intProperty;
        StringProperty = stringProperty;
        Values = values;
    }
}
```

The Actor System

- Manages actor life cycles
- Manages messaging
- Manages inboxes
- Manages thread scheduling
- Manages the system event bus
- ...



The Actor System

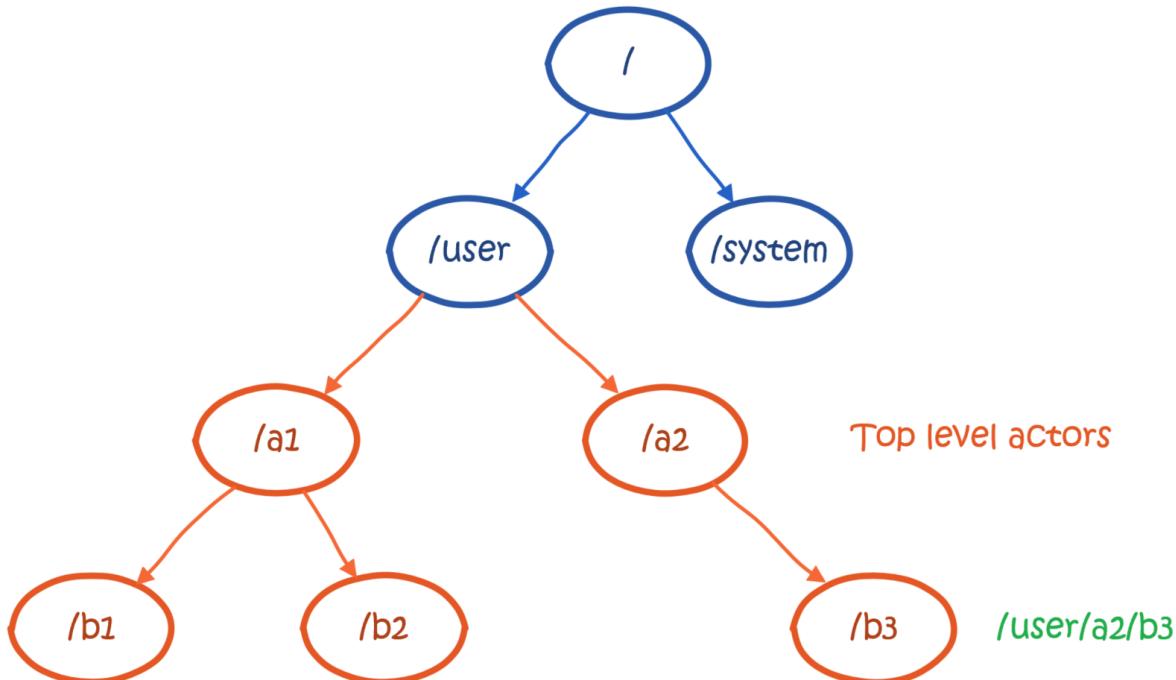
```
public static void Main(string[] args)
{
    var actorSystem = ActorSystem.Create("MyActorSystem");

    var myActorProps = Props.Create<MyActor>();

    // Creates /user/my-actor-name
    var myActor = actorSystem.ActorOf(myActorProps, "my-actor-name");
    IActorRef
    var message = new MyMessage(1, "Hello", ImmutableArray<decimal>.Empty);

    myActor.Tell(message);
}
```

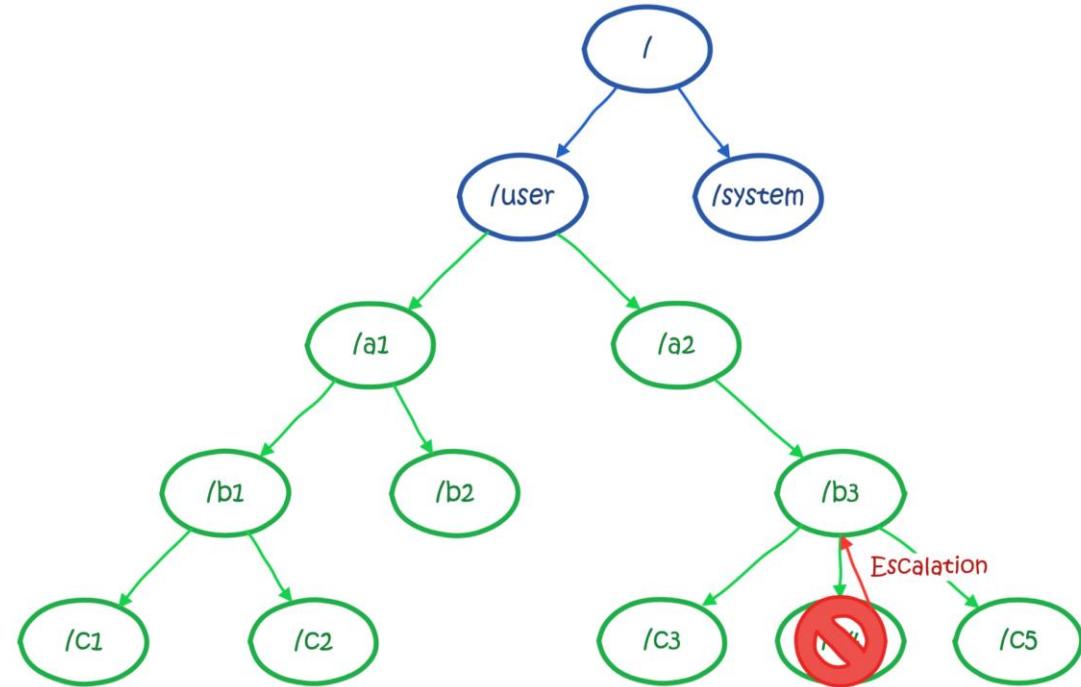
The actor hierarchy



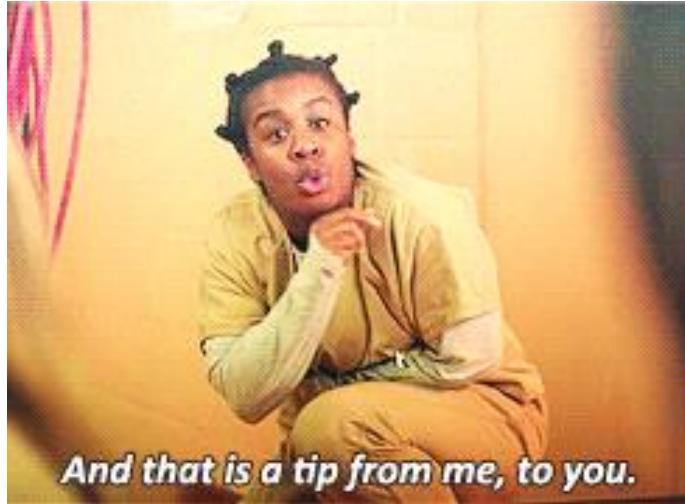
- Actors can have children
- Position in the hierarchy = address
- Directly under /user = 'top level actors'
- 3 default actors:
 - /
 - /user
 - /system

Fault Tolerance - Supervision

- Errors are escalated to the parent
- Parent decides OR escalates further
- Strategy:
 - OneForOne: only the failing actor
 - OneForAll: all children
- Action:
 - Resume
 - Stop
 - Restart



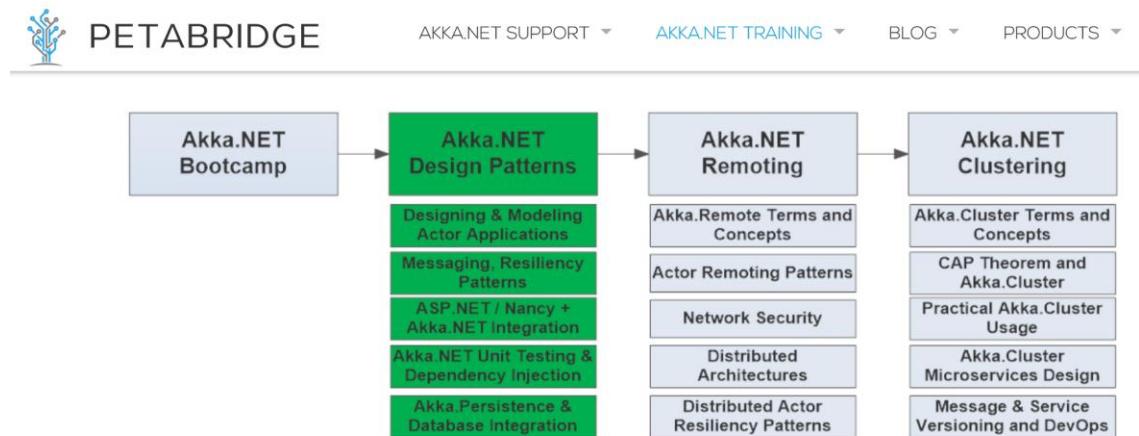
General development ideas



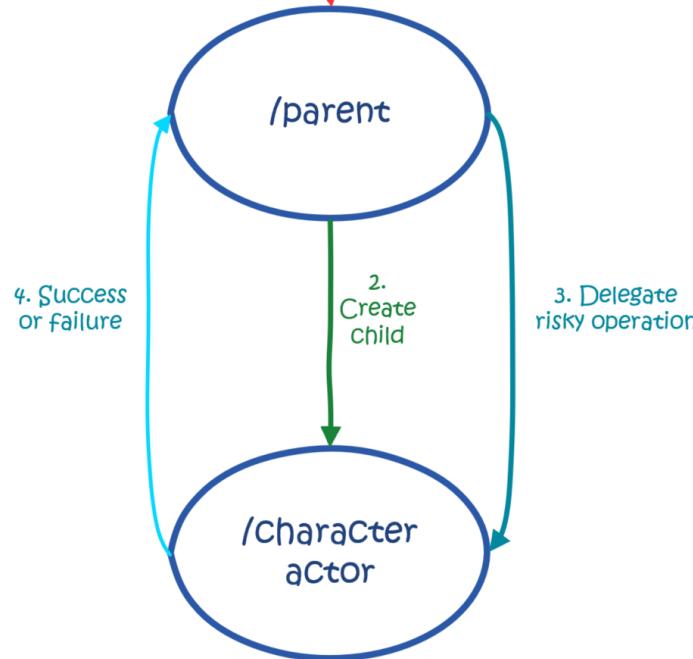
- Split workloads into small chunks
- Make separate actors for every task
- Push risk to the edges,
handle faults there
- Avoid ‘bottleneck actors’

New design patterns

- Fan-out Pattern
- Parent Proxy Pattern
- Consensus Pattern
- Character Actor
- ...



1. Request to do risky operation



Character Actor



The problem domain

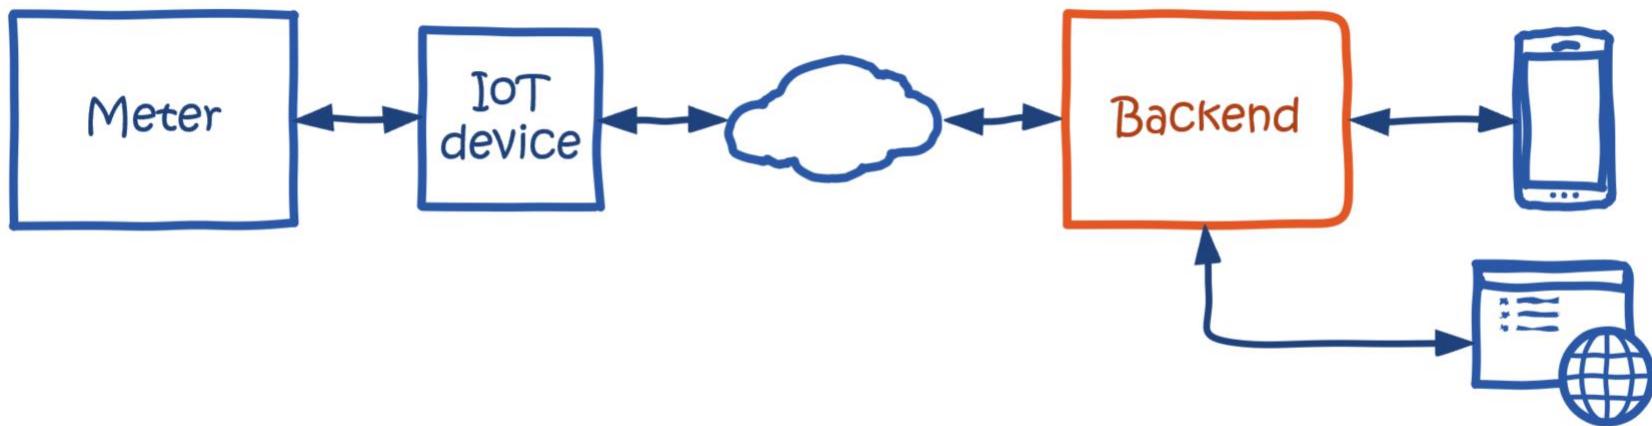
What are we going to solve?

IT IS ABOUT PEOPLE

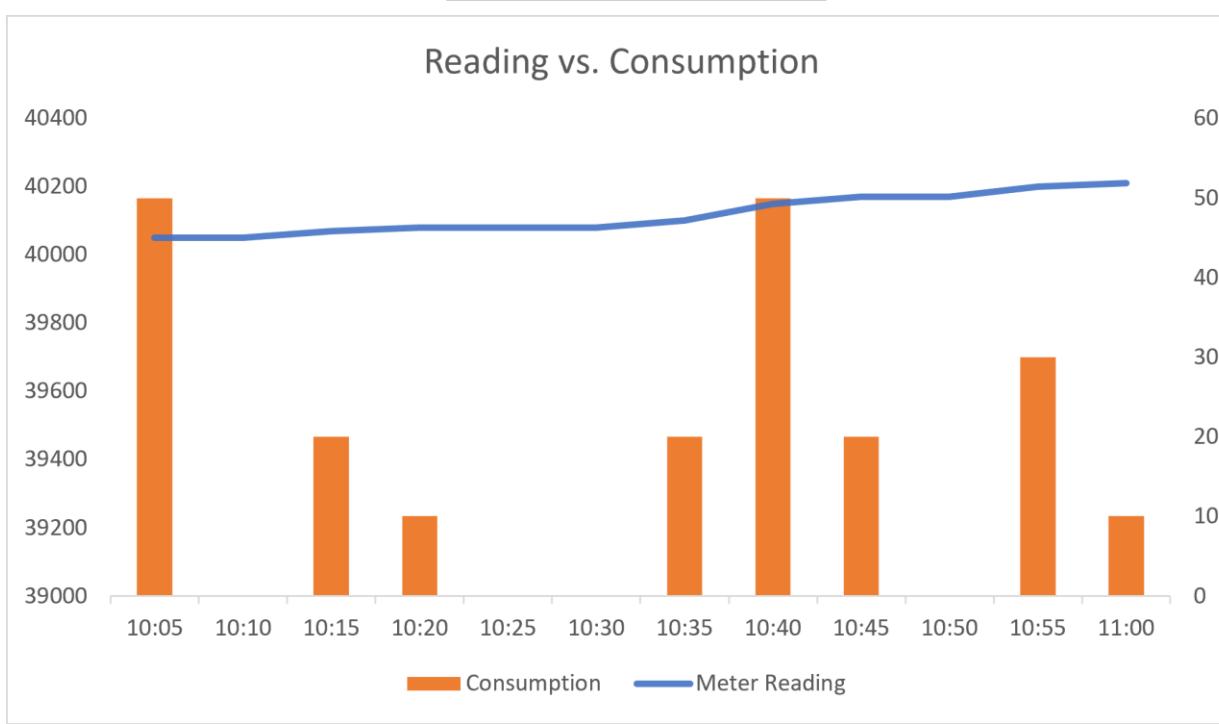




Connection situation



Reading vs Consumption



What are we interested in?

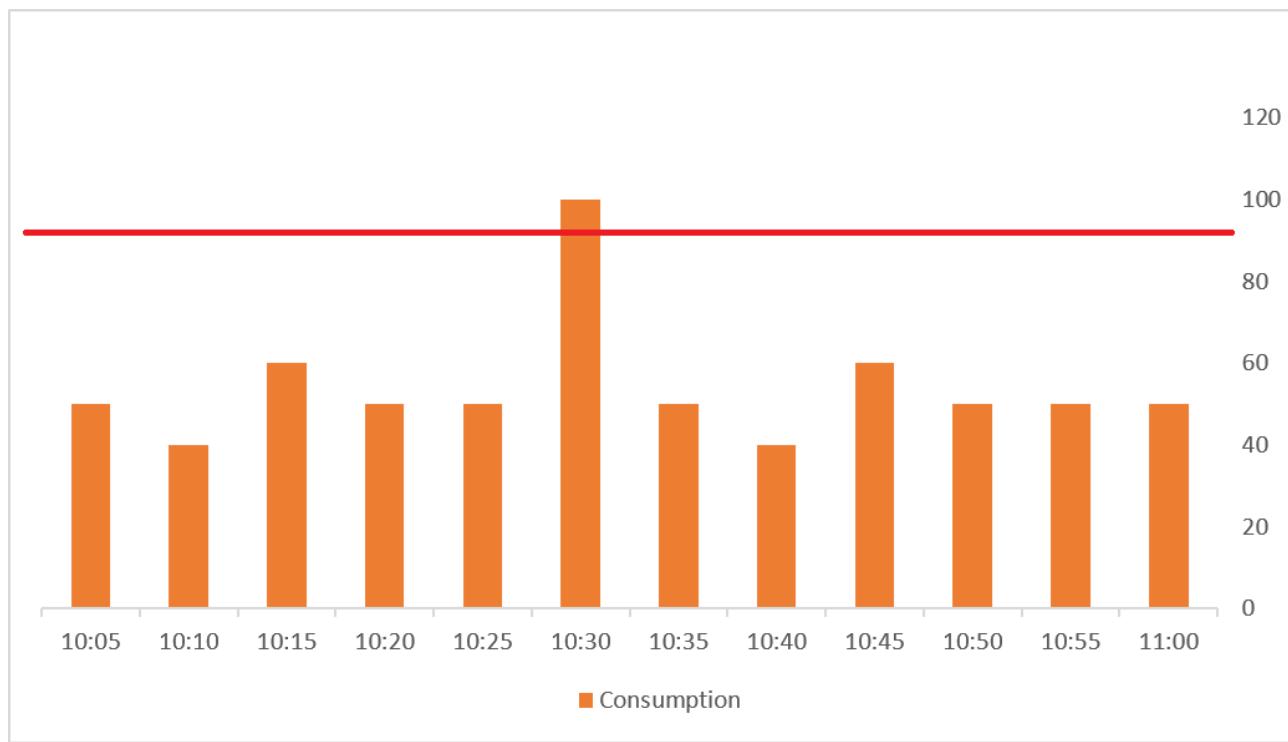
Storage of historic usage

- Storage of (normalized) values
- Plotting of consumption graphs
- Comparison of time periods

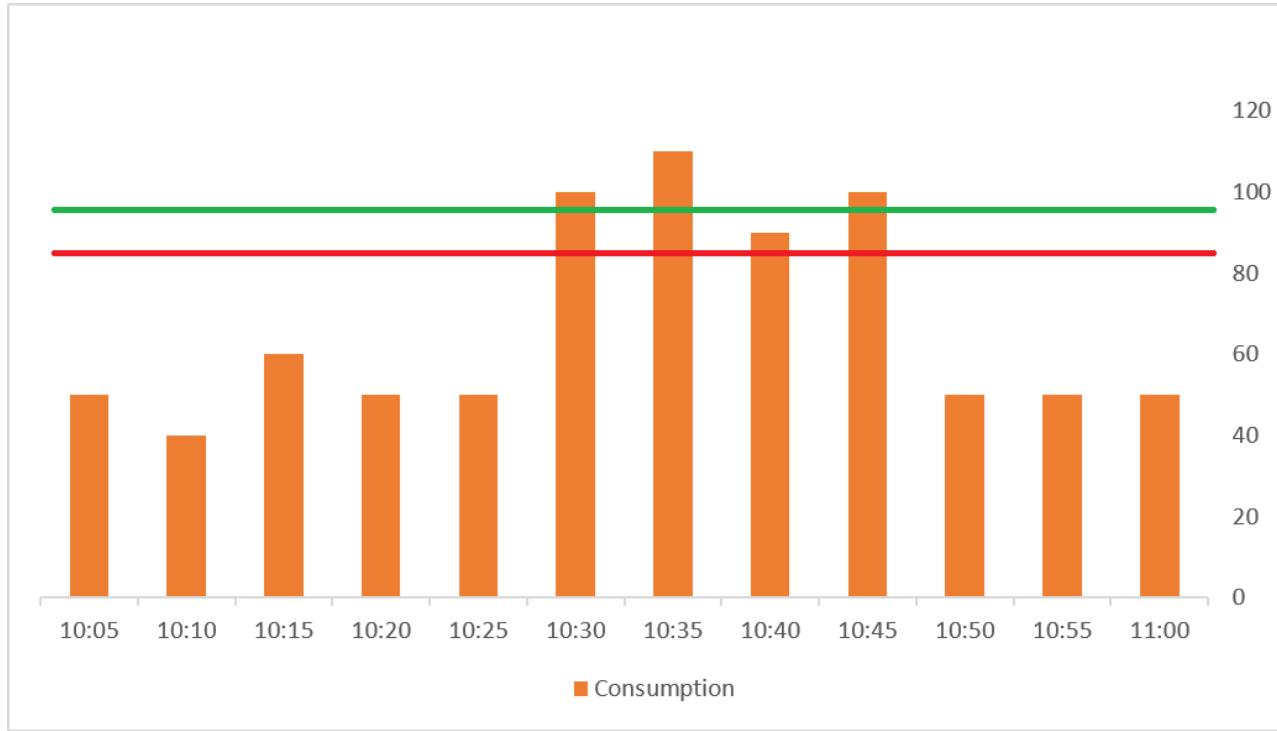
Alerting

- When momentary consumption exceeds threshold
- When a threshold is exceeded for a period

Momentary threshold



Periodic threshold



Periodic threshold

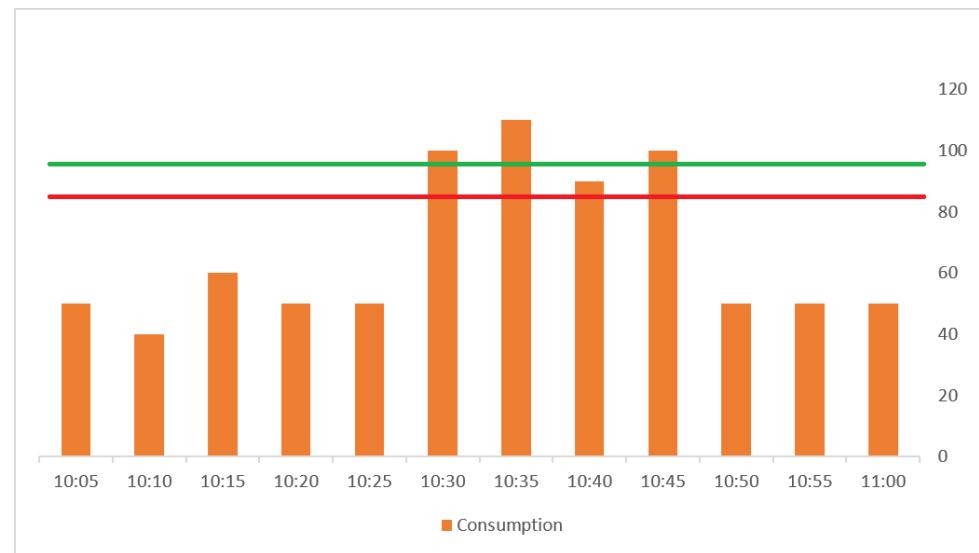
Alert after 20 minutes

Green line:

- Exceeded on average
- Not exceeded every time block

Red line:

- Exceeded on average
- Exceeded every time block



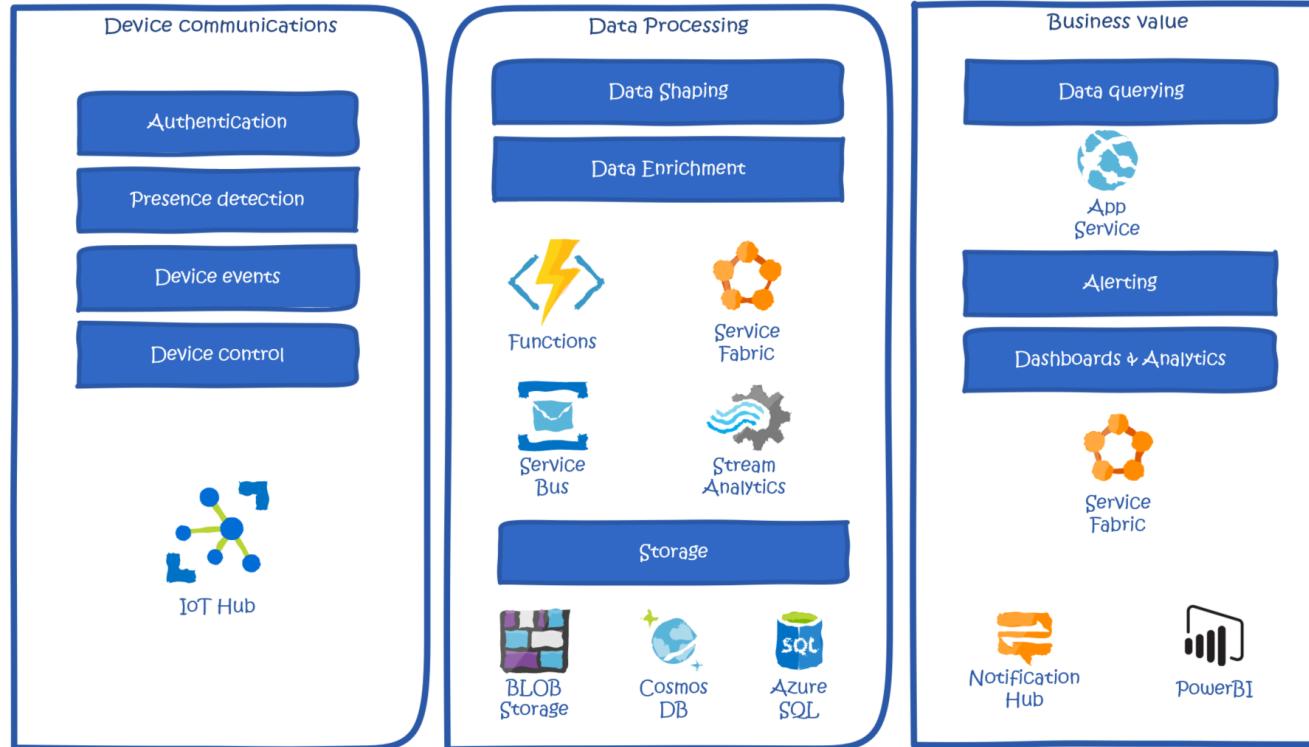
How does Akka.NET fit?

What part of the solution can Akka.NET provide?

IT IS ABOUT PEOPLE

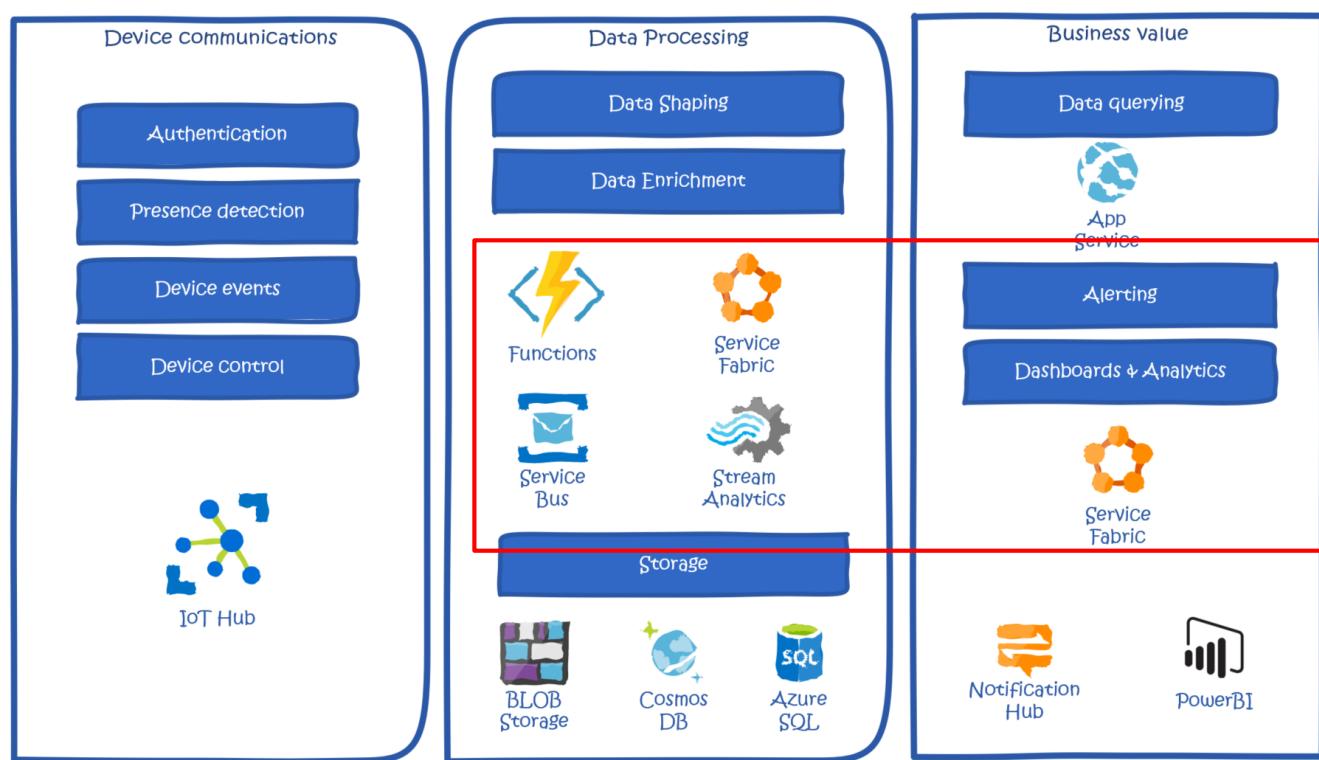


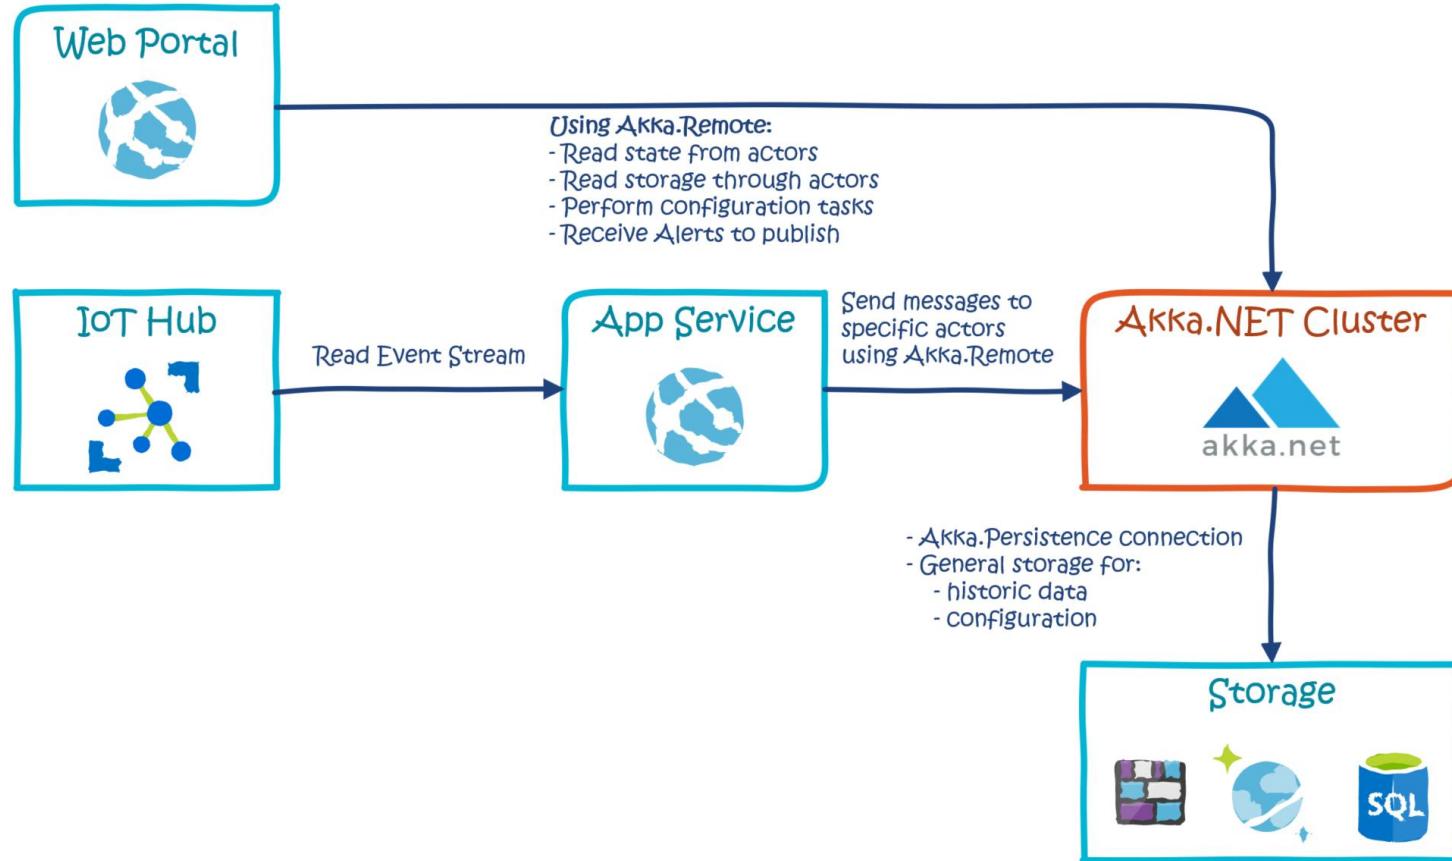
Your typical IoT stack*



* With Microsoft Azure service examples

What can we use Akka.NET for?





Don't be a magpie!

Good fits:

- Gaming backends
- Trading systems
- Internet of Things
- Parallelizable calculations
- ... any stateful high throughput application

It doesn't have to be the whole solution!



Implementation details

How can you use Akka.NET in this scenario?

IT IS ABOUT PEOPLE



IT IS ABOUT PEOPLE



- 1 Getting messages to the ActorSystem
- 2 Normalizing measurements
- 3 Persisting Data
- 4 Restart behavior

Getting messages to the ActorSystem

How to use Remoting & Proxy actors

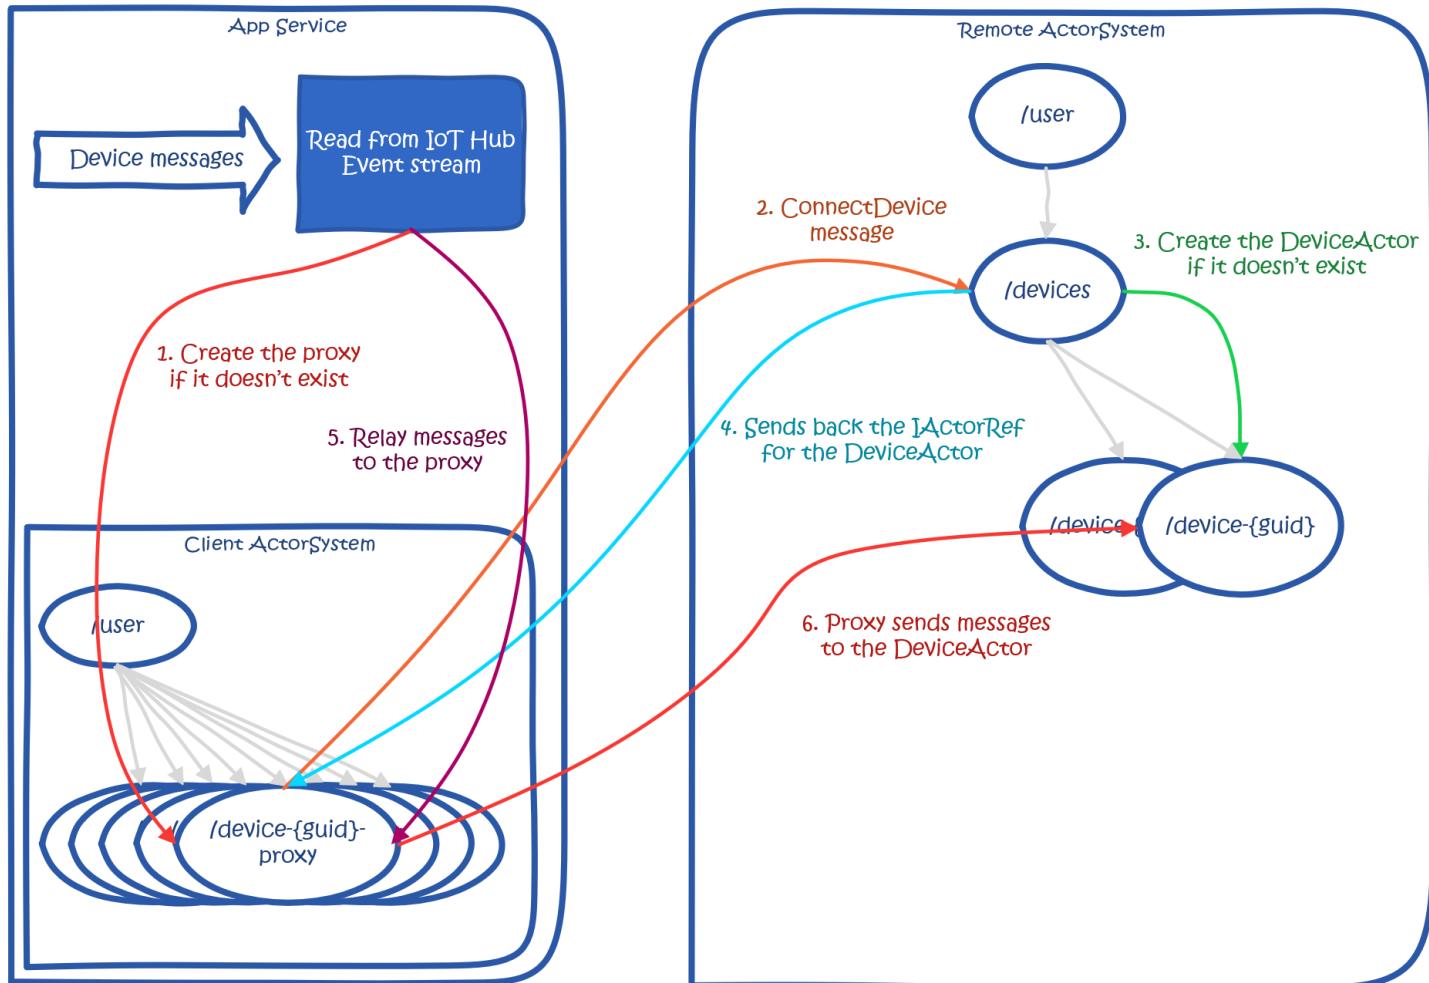
Akka.Remote

ActorSystems can talk to other ActorSystems

- Remote addressing
- Remote deployment
- Remote messaging
- Location Transparency
- Multiple transports

All parts form an "ActorPath"





DeviceActorProxy

```
class DeviceActorProxy : ReceiveActor
{
    private readonly Guid _deviceId;
    private IActorRef _deviceActor;

    public DeviceActorProxy(Guid deviceId) [...]
```

protected override void PreStart() Runs BEFORE the Actor starts to accept messages.

```
{
```

var devicesActorPath = \$"{Constants.RemoteActorSystemAddress}/user/devices";
 var devicesActor = Context.ActorSelection(devicesActorPath);

IActorSelection Creates an ActorSelection to the DevicesActor in the REMOTE ActorSystem

```
    var request = new ConnectDevice(_deviceId);  
    devicesActor.Tell(request); Sends the 'connect' message to the DevicesActor
```

```
}
```

private void HandleDeviceConnected(DeviceConnected message)

```
{
    _deviceActor = message.DeviceRef; When the DevicesActor responds, we will have the IActorRef to the DeviceActor
```

```
}
```

DeviceActorProxy

```
class DeviceActorProxy : ReceiveActor
{
    private readonly Guid _deviceId;
    private IActorRef _deviceActor;

    public DeviceActorProxy(Guid deviceId) More robust than an UntypedActor
    {
        _deviceId = deviceId;
        Receive<MeterReadingReceived>(HandleMeterReadingReceived);
        Receive<DeviceConnected>(HandleDeviceConnected); This is how you configure message types
    }

    protected override void PreStart()...

    private void HandleDeviceConnected(DeviceConnected message)...

    private void HandleMeterReadingReceived(MeterReadingReceived message)
    {
        _deviceActor?.Tell(message); Just forward to the actual DeviceActor
    }

    public static Props CreateProps(Guid deviceId) This makes refactoring SO MUCH easier
    {
        return Props.Create<DeviceActorProxy>(deviceId);
    }
}
```

DevicesActor

```
private void HandleConnectDevice(ConnectDevice request)
{
    if (!_deviceActors.ContainsKey(request.Id))           Check if the device has already been created
    {
        CreateDeviceActor(request.Id);
    }
    var response = new DeviceConnected(_deviceActors[request.Id]);
    Sender.Tell(response);                            Sender holds the ActorRef to the sender of the current message
}

private void CreateDeviceActor(Guid deviceId)
{
    var props = DeviceActor.CreateProps(deviceId);
    var name = $"device-{deviceId}";
    var deviceActorRef = Context.ActorOf(props, name);      Addresses need to be unique.
                                                               Use the Context to create child Actors.

    _deviceActors[deviceId] = deviceActorRef;
}
```

Normalizing measurements

Making sure downstream actors get consistent data

Why Normalization?

Writing business logic is easier when you know values are consistent:

- Consistent timestamps
- No gaps
- Incorrect values filtered
- ...

Deal with it in one place



Timestamp correction & buckets

RAW			NORMALIZED		
Timestamp	Reading	Consumption	Timestamp	Reading	Consumption
9:59:25	40000				
10:00:25	40060	60	10:00:00	40035	
10:01:25	40120	60			
10:02:25	40180	60			
10:03:25	40240	60			
10:04:25	40300	60			
10:05:25	40360	60	10:05:00	40335	300
10:06:25	40420	60			
10:07:25	40480	60			
10:08:25	40540	60			
10:09:25	40600	60			
10:10:25	40660	60	10:10:00	40635	300
10:11:25	40720	60			
10:12:25	40780	60			
10:13:25	40840	60			
10:14:25	40900	60			
10:15:25	40960	60	10:15:00	40935	300

Gap filling

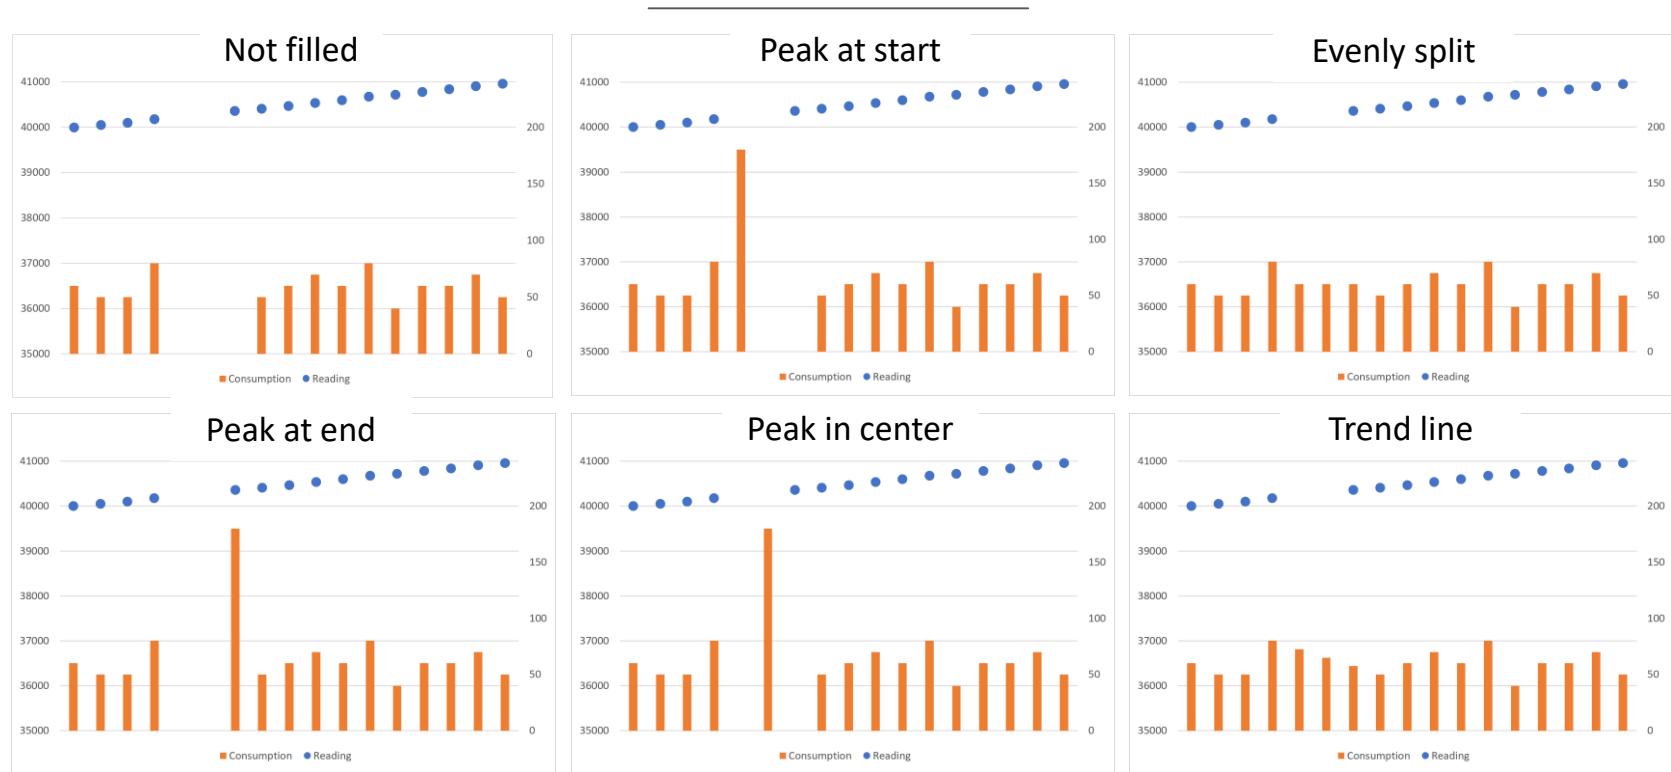


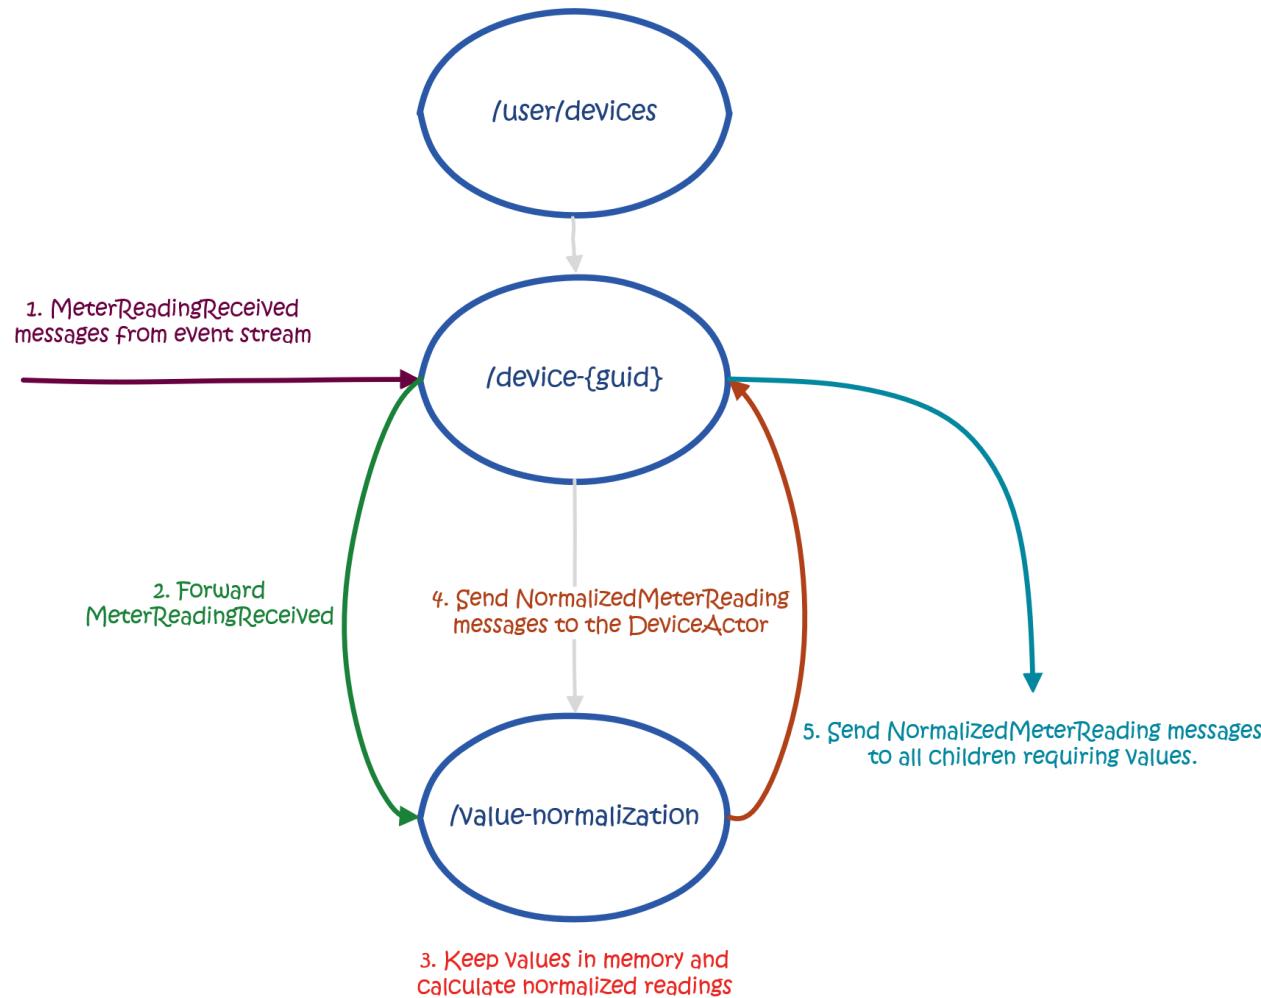
- Do we want to fill this gap?
 - If so, how?
 - Do other Actors need to know?
If yes, add a flag to the message

There is no ‘right’ answer

Just one that fits your case

Possible filling solutions





Persisting Data

Saving to disk what cannot be lost

Akka.Persistence

Actors that recover their state when created:

- Inherit from `PersistentActor`
- Give it a unique `PersistenceId`
- Persist Events with the `Persist(...)` command
- Persist snapshots with the `SaveSnapshot(...)` command
- Register `Recover<T>(...)` handlers to restore state

```
public class MyPersistedActor : ReceivePersistentActor
{
    // Any PersistentActor needs a unique key!
    public override string PersistenceId { get; }

    // Grouping state into a state object is a good idea
    private MyState _state = new MyState();

    public MyPersistedActor(Guid id)
    {
        PersistenceId = $"my-persisted-actor-{id}";

        // There's a difference between 'Commands' and 'Recovers'
        Command<MyMessage>(HandleCommand);
        Recover<MyMessage>(HandleMessageInternal);

        // Snapshot events
        Recover<SnapshotOffer>(HandleSnapshotOffer);
        Command<SaveSnapshotSuccess>(HandleSnapshotSuccess);
        Command<SaveSnapshotFailure>(HandleSnapshotFailure);
    }
}
```

```
private int _msgSinceLastSnapshot = 0;

private void HandleCommand(MyMessage command)
{
    // Persists the message to the store and the actor simultaneously.
    Persist<MyMessage>(command, HandleMessageInternal);

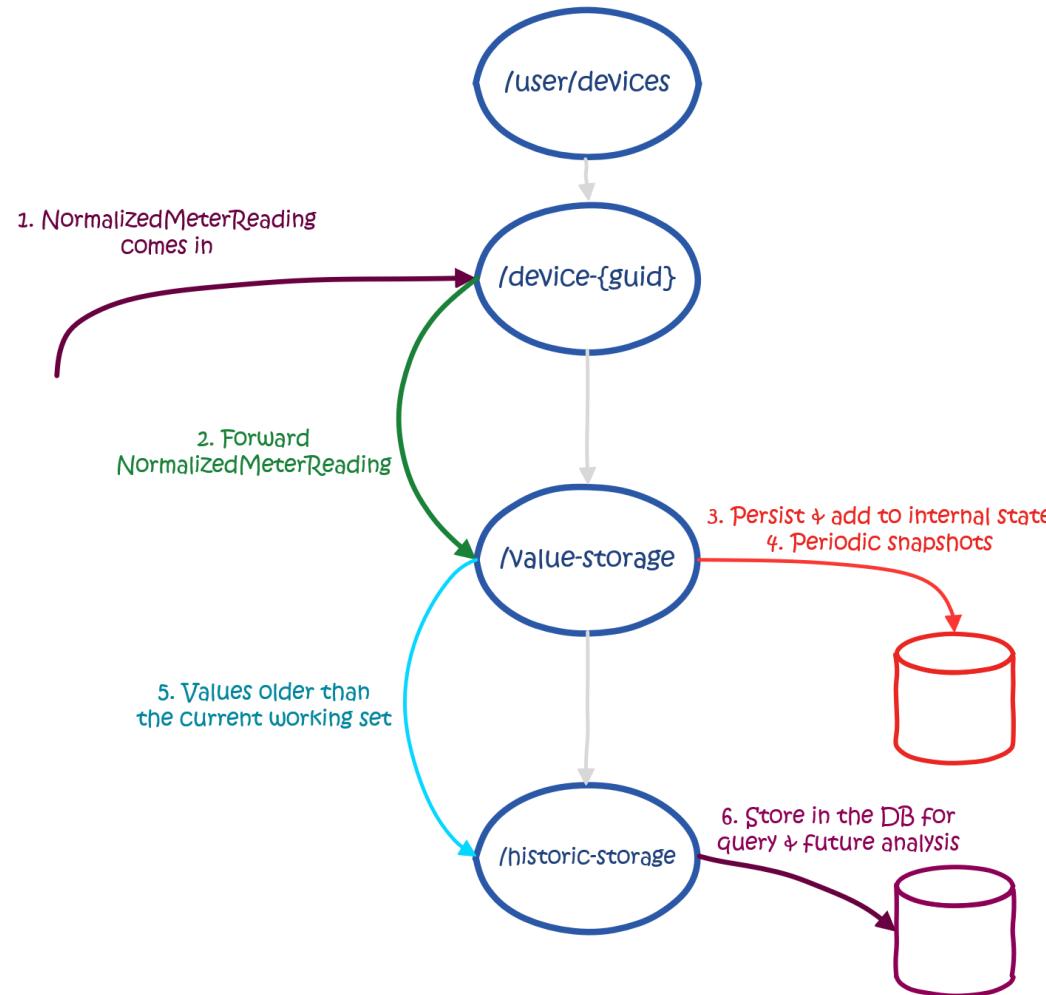
    // Save a snapshot every 100 messages
    if (_msgSinceLastSnapshot == 100)
    {
        SaveSnapshot(_state);
        _msgSinceLastSnapshot = 0;
    }
}

private void HandleMessageInternal(MyMessage message)
{
    // In recovery, we call this directly, no need to persist it again.
    _state.Add(message);
    _msgSinceLastSnapshot++;
}
```

```
private void HandleSnapshotOffer(SnapshotOffer offer)
{
    if (offer.Snapshot is MyState newState)
        _state = newState;
}

private void HandleSnapshotSuccess(SaveSnapshotSuccess success)
{
    // Handle a successful snapshot save
}

private void HandleSnapshotFailure(SaveSnapshotFailure failure)
{
    // Handle the failure to save a snapshot
}
```



Restart Behavior

How can we be sure to get going again after a restart?

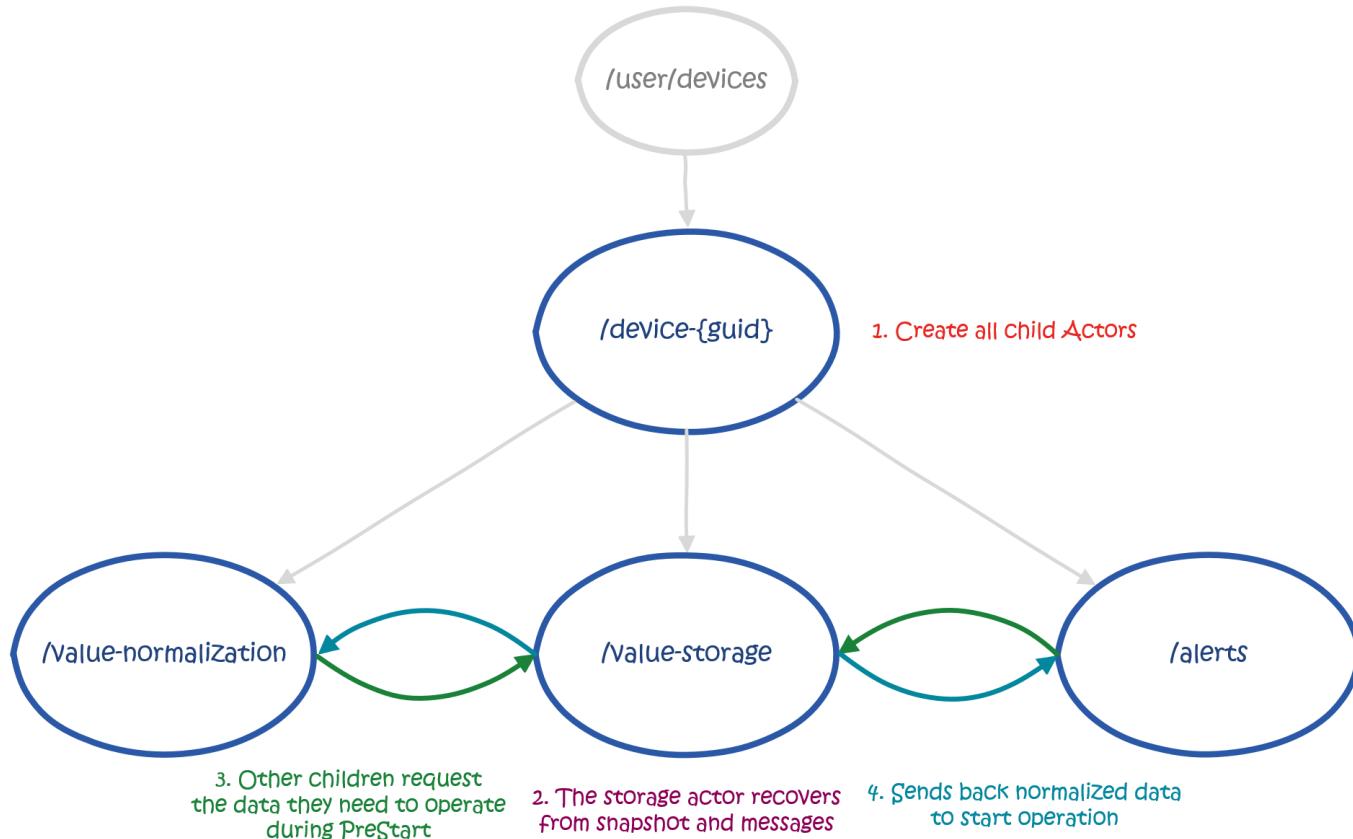
After a system restart

Recreating Actors isn't that hard:

- Query the DB on startup
- Create the required Actors

How do we get all our Actor state back?

- Minimize the amount of actors that need to recover state
- 1 PersistedActor per device = ideal
- Let other actors query that actor for the state they need



Beyond this talk ...

The stuff that we didn't talk about ...

IT IS ABOUT PEOPLE



Make Akka.NET production ready

- **Configuration:**
HOCON
- **Clustering:**
Run across multiple machines
- **Logging:**
Adapters for Nlog, SeriLog, etc.
- **Dependency Injection:**
Akka.NET supports DI for your actors
- **Production monitoring:**
Phobos



Start learning



- **FREE Akka.NET Bootcamp by Petabridge:**
<https://github.com/petabridge/akka-bootcamp>
- **PluralSight courses:**
There are some good courses on there!
- **Petabridge blog:**
<https://petabridge.com/blog/>
- **Petabridge remote training (paid):**
Worth it when you have serious questions

Deployment

1. Pause the process that reads from the event stream
2. Wait for processing to end
3. Deploy the Akka.NET cluster
4. Re-create actors (triggering Persistence restores)
5. Resume sending from the event stream

→ When done right, you can do this without losing data!

Conclusion

1. Check if your problem domain is a fit for Actors
2. Decide which part of the solution will be Akka.NET
3. Design your actor hierarchies appropriately
4. Normalizing data helps a lot
5. Think about deployment & recycles

Ask me some questions!



About

Hannes Lowette

.NET Consultant & Competence Coach at **@Axxes_IT**



@hannes_lowette



#20086521



Code samples and slides at :

github.com/Belenar/Axxes.AkkaDotNet.SensorData

Pictures from:

<https://thetravelquandary.com/>

<https://imgflip.com/>

<https://commons.wikimedia.org/>

<https://akka.io/>

<https://codingjourneyman.com/>

<https://getakka.net/>

<https://dotnet.github.io/orleans/>

<https://blogs.msdn.microsoft.com/>

<https://www.karlrupp.net/>

<https://www.xkcd.com/>

<http://www.last.fm/>

<https://www.newscientist.com/>

<https://io9.gizmodo.com/>

<https://www.thebestcolleges.org/>